

uThreads

0.3.0

Generated by Doxygen 1.8.11

Contents

1	uThreads: Concurrent User Threads in C++(and C)	1
2	Hierarchical Index	7
2.1	Class Hierarchy	7
3	Class Index	9
3.1	Class List	9
4	File Index	11
4.1	File List	11
5	Class Documentation	13
5.1	BlockingQueue Class Reference	13
5.1.1	Detailed Description	14
5.1.2	Member Function Documentation	14
5.1.2.1	signal(std::mutex &lock, uThread *&owner)	14
5.1.2.2	signal(std::mutex &lock)	14
5.1.2.3	signal(Mutex &)	14
5.1.2.4	signalAll(Mutex &)	15
5.1.2.5	suspend(std::mutex &lock)	15
5.1.2.6	suspend(Mutex &)	15
5.2	Cluster Class Reference	16
5.2.1	Detailed Description	16
5.2.2	Constructor & Destructor Documentation	17
5.2.2.1	Cluster()	17

5.2.2.2	Cluster(const Cluster &)=delete	17
5.2.3	Member Function Documentation	17
5.2.3.1	getDefaultCluster()	17
5.2.3.2	getID() const	17
5.2.3.3	getNumberOfThreads() const	17
5.2.3.4	operator=(const Cluster &)=delete	17
5.3	ConditionVariable Class Reference	18
5.3.1	Detailed Description	18
5.3.2	Member Function Documentation	18
5.3.2.1	empty()	18
5.3.2.2	signal(Mutex &mutex)	18
5.3.2.3	signalAll(Mutex &mutex)	18
5.3.2.4	wait(Mutex &mutex)	19
5.4	Connection Class Reference	19
5.4.1	Detailed Description	20
5.4.2	Constructor & Destructor Documentation	20
5.4.2.1	Connection()	20
5.4.2.2	Connection(int fd)	20
5.4.2.3	Connection(int domain, int type, int protocol)	21
5.4.3	Member Function Documentation	21
5.4.3.1	accept(Connection *conn, struct sockaddr *addr, socklen_t *addrlen)	21
5.4.3.2	accept(struct sockaddr *addr, socklen_t *addrlen)	21
5.4.3.3	bind(const struct sockaddr *addr, socklen_t addrlen)	21
5.4.3.4	close()	22
5.4.3.5	connect(const struct sockaddr *addr, socklen_t addrlen)	22
5.4.3.6	getFd() const	22
5.4.3.7	listen(int backlog)	22
5.4.3.8	read(void *buf, size_t count)	22
5.4.3.9	recv(void *buf, size_t len, int flags)	23
5.4.3.10	recvfrom(void *buf, size_t len, int flags, struct sockaddr *src_addr, socklen_t *addrlen)	23

5.4.3.11	recvmsg(int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags, struct timespec *timeout)	23
5.4.3.12	recvmsg(int sockfd, struct msghdr *msg, int flags)	24
5.4.3.13	send(const void *buf, size_t len, int flags)	24
5.4.3.14	sendmmsg(int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags)	24
5.4.3.15	sendmsg(const struct msghdr *msg, int flags)	25
5.4.3.16	sendto(int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen)	25
5.4.3.17	socket(int domain, int type, int protocol)	25
5.4.3.18	write(const void *buf, size_t count)	26
5.5	IOPoller Class Reference	26
5.6	kThread Class Reference	26
5.6.1	Detailed Description	28
5.6.2	Constructor & Destructor Documentation	28
5.6.2.1	kThread(Cluster &)	28
5.6.3	Member Function Documentation	28
5.6.3.1	currentkThread()	28
5.6.3.2	getID()	28
5.6.3.3	getThreadNativeHandle()	29
5.6.3.4	getTotalNumberOfkThreads()	29
5.6.3.5	operator=(const kThread &)=delete	29
5.7	Mutex Class Reference	29
5.7.1	Detailed Description	31
5.7.2	Member Function Documentation	31
5.7.2.1	acquire()	31
5.8	OwnerLock Class Reference	31
5.8.1	Detailed Description	32
5.8.2	Member Function Documentation	32
5.8.2.1	acquire()	32
5.8.2.2	release()	33
5.9	Semaphore Class Reference	33

5.9.1	Detailed Description	33
5.9.2	Constructor & Destructor Documentation	33
5.9.2.1	Semaphore(mword c=0)	33
5.9.3	Member Function Documentation	34
5.9.3.1	P()	34
5.10	uThread Class Reference	34
5.10.1	Detailed Description	36
5.10.2	Constructor & Destructor Documentation	37
5.10.2.1	uThread(vaddr sb, size_t ss)	37
5.10.3	Member Function Documentation	37
5.10.3.1	create(size_t ss, bool joinable=false)	37
5.10.3.2	create(bool joinable=false)	37
5.10.3.3	currentUThread()	37
5.10.3.4	getCurrentCluster() const	38
5.10.3.5	getID() const	38
5.10.3.6	getTotalNumberOfUTs()	38
5.10.3.7	join()	38
5.10.3.8	migrate(Cluster *)	38
5.10.3.9	operator=(const uThread &)=delete	38
5.10.3.10	start(const Cluster &cluster, ptr_t func, ptr_t arg1=nullptr, ptr_t arg2=nullptr, ptr_t arg3=nullptr)	39
5.10.3.11	terminate()	39
5.10.3.12	yield()	39
5.11	uThreadCache Class Reference	39
5.11.1	Detailed Description	40
5.11.2	Member Function Documentation	40
5.11.2.1	pop()	40
5.11.2.2	push(uThread *ut)	40

6	File Documentation	41
6.1	/home/saman/Programming/Research/uThreads/src/cwrapper.h File Reference	41
6.1.1	Detailed Description	45
6.1.2	Function Documentation	45
6.1.2.1	cluster_create()	45
6.1.2.2	cluster_destroy(WCluster *cluster)	45
6.1.2.3	cluster_get_current()	45
6.1.2.4	cluster_get_default()	45
6.1.2.5	cluster_get_id(WCluster *cluster)	46
6.1.2.6	cluster_get_number_of_kThreads(WCluster *cluster)	46
6.1.2.7	condition_variable_create()	46
6.1.2.8	condition_variable_destroy(WConditionVariable *cv)	46
6.1.2.9	condition_variable_empty(WConditionVariable *cv)	46
6.1.2.10	condition_variable_signal(WConditionVariable *cv, WMutex *mutex)	46
6.1.2.11	condition_variable_signal_all(WConditionVariable *cv, WMutex *mutex)	46
6.1.2.12	condition_variable_wait(WConditionVariable *cv, WMutex *mutex)	47
6.1.2.13	connection_accept(WConnection *acceptor, WConnection *conn, struct sockaddr *addr, socklen_t *addrlen)	47
6.1.2.14	connection_accept_connection(WConnection *acceptor, struct sockaddr *addr, socklen_t *addrlen)	47
6.1.2.15	connection_bind(WConnection *conn, const struct sockaddr *addr, socklen_t addrlen)	47
6.1.2.16	connection_block_on_read(WConnection *conn)	47
6.1.2.17	connection_block_on_write(WConnection *conn)	48
6.1.2.18	connection_close(WConnection *conn)	48
6.1.2.19	connection_connect(WConnection *conn, const struct sockaddr *addr, socklen_t addrlen)	48
6.1.2.20	connection_create()	48
6.1.2.21	connection_create_socket(int domain, int type, int protocol)	48
6.1.2.22	connection_create_with_fd(int fd)	48
6.1.2.23	connection_destroy(WConnection *c)	49
6.1.2.24	connection_get_fd(WConnection *conn)	49

6.1.2.25	<code>connection_listen(WConnection *conn, int backlog)</code>	49
6.1.2.26	<code>connection_read(WConnection *conn, void *buf, size_t count)</code>	49
6.1.2.27	<code>connection_rcv(WConnection *conn, void *buf, size_t len, int flags)</code>	49
6.1.2.28	<code>connection_rcvfrom(WConnection *conn, void *buf, size_t len, int flags, struct sockaddr *src_addr, socklen_t *addrlen)</code>	50
6.1.2.29	<code>connection_rcvmsg(WConnection *conn, int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags, struct timespec *timeout)</code>	50
6.1.2.30	<code>connection_rcvmsg(WConnection *conn, int sockfd, struct msghdr *msg, int flags)</code>	50
6.1.2.31	<code>connection_send(WConnection *conn, const void *buf, size_t len, int flags)</code>	51
6.1.2.32	<code>connection_sendmsg(WConnection *conn, int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags)</code>	51
6.1.2.33	<code>connection_sendmsg(WConnection *conn, const struct msghdr *msg, int flags)</code>	51
6.1.2.34	<code>connection_sendto(WConnection *conn, int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen)</code>	52
6.1.2.35	<code>connection_socket(WConnection *conn, int domain, int type, int protocol)</code>	52
6.1.2.36	<code>connection_write(WConnection *conn, const void *buf, size_t count)</code>	52
6.1.2.37	<code>kThread_create(WCluster *cluster)</code>	53
6.1.2.38	<code>kThread_destroy(WkThread *kt)</code>	53
6.1.2.39	<code>kThread_get_current()</code>	53
6.1.2.40	<code>kThread_get_current_pthread_id()</code>	53
6.1.2.41	<code>kThread_get_pthread_id(WkThread *kt)</code>	53
6.1.2.42	<code>kThread_get_total_number_of_kThreads()</code>	53
6.1.2.43	<code>mutex_acquire(WMutex *mutex)</code>	53
6.1.2.44	<code>mutex_create()</code>	54
6.1.2.45	<code>mutex_destroy(WMutex *mutex)</code>	54
6.1.2.46	<code>mutex_release(WMutex *mutex)</code>	54
6.1.2.47	<code>ownerlock_acquire(WOwnerLock *olock)</code>	54
6.1.2.48	<code>ownerlock_create()</code>	54
6.1.2.49	<code>ownerlock_destroy(WOwnerLock *olock)</code>	54
6.1.2.50	<code>ownerlock_release(WOwnerLock *olock)</code>	54
6.1.2.51	<code>semaphore_create()</code>	54
6.1.2.52	<code>semaphore_destroy(WSemaphore *sem)</code>	54

6.1.2.53	<code>semaphore_p(WSemaphore *sem)</code>	54
6.1.2.54	<code>semaphore_v(WSemaphore *sem)</code>	54
6.1.2.55	<code>uThread_create(bool joinable)</code>	54
6.1.2.56	<code>uThread_detach(WuThread *ut)</code>	55
6.1.2.57	<code>uThread_get_current()</code>	55
6.1.2.58	<code>uThread_get_id(WuThread *ut)</code>	55
6.1.2.59	<code>uThread_get_total_number_of_uThreads()</code>	55
6.1.2.60	<code>uThread_join(WuThread *ut)</code>	55
6.1.2.61	<code>uThread_migrate(WCluster *cluster)</code>	55
6.1.2.62	<code>uThread_start(WuThread *ut, WCluster *cluster, void *func, void *arg1, void *arg2, void *arg3)</code>	56
6.1.2.63	<code>uThread_terminate(WuThread *ut)</code>	56
6.1.2.64	<code>uThread_yield()</code>	56

Chapter 1

uThreads: Concurrent User Threads in C++(and C)

What are uThreads?

uThreads is a concurrent library based on cooperative scheduling of user-level threads(fibers) implemented in C++. User-level threads are lightweight threads that execute on top of kernel threads to provide concurrency as well as parallelism. Kernel threads are necessary to utilize processors, but they come with the following drawbacks:

- Each suspend/resume operation involves a kernel context switch
- Thread preemption causes additional overhead
- Thread priorities and advanced scheduling causes additional overhead

Cooperative user-level threads, on the other hand, provide light weight context switches and omit the additional overhead of preemption and kernel scheduling. Most Operating Systems only support a 1:1 thread mapping (1 user-level thread to 1 kernel-level thread), where multiple kernel threads execute at the same time to utilize multiple cores and provide parallelism. e.g., Linux supports only 1:1 thread mapping. There is also N:1 thread mapping, where multiple user-level threads can be mapped to a single kernel-level thread. The kernel thread is not aware of the user-level threads existence. For example, Facebook's `folly::fiber`, `libmill`, and `libtask` use N:1 mapping. Having N:1 mapping means if the application blocks at the kernel level, all user-level threads are blocked and application cannot move forward. One way to address this is to only block on user level, hence, blocking user-level threads. This setting works very well with IO bound applications, however, if a user thread requires using a CPU for a while, it can block other user threads and the task is better to be executed asynchronously on another core to prevent this from happening. In order to avoid this problem, user threads can be mapped to multiple kernel-level threads. Thus, creating the third scenario with M:N or hybrid mapping. e.g., `go` and `uC++` use M:N mapping.

uThreads supports M:N mapping of *uThreads* (user-level threads) over *kThreads* (kernel-level threads) with cooperative scheduling. *kThreads* can be grouped together by *Clusters*, and uThreads can migrate among Clusters. Figure 1 shows the structure of an application implemented using uThreads using a single ReadyQueue Scheduler. You can find the documentation here <http://samanbarghi.github.io/uThreads>

Webserver throughput results

Here are a comparison of a simple webserver throughput with `fasthttp`, `cppsp`, and `nodejs`. Experiments are executed on my laptop (i7 quad core). Note that not much optimization is applied to any of the applications, thus there might be some space to squeeze more throughput out of these applications. You can check the source code of the sample webserver under the test directory. All servers return a "hello world" response, and the client (in this case wrk) throws a huge number of concurrent and pipelined requests at each server. This experiment shows the overhead of each framework since the response is kept very small (similar to TechEmpower "plaintext benchmark").

Dependencies

Currently uThreads only supports Linux on x86_64 platforms. It also depends on the following:

- gcc > 4.8
- linux kernel >= 2.9

Building and installation

To build and install the library issue the following commands:

```
1 make
2 make install
```

Currently, everything is installed under `/usr/local/lib` and `/usr/local/include/uThreads`. To change this you need to edit the Makefile and change the `DEST_DIR`.

Usage

- Include "uThreads/uThreads.h" in your source file.
- Link your program with uThreads library (-luThreads) at compile time.

There are sample applications under `test` directory, to compile them issue `make test`, and you can find the binaries under `bin` directory. Refer to the [documentation](#) for more information.

uThreads structure

This section explains the internals of uThreads.

Basics

Clusters are used to group kThreads together. Each **Cluster** can contain one or more kThreads, but each **kThread** only belongs to a single **Cluster**. Each **Cluster** includes a *Scheduler* which is used to schedule uThreads over kThreads in that **Cluster**. Application programmer decides how many kThreads belong to a **Cluster** by assigning them upon creation. *Clusters* can be used to execute different tasks over separate kThreads and if pinned properly, over separate cores. For example they can be used to provide better CPU cache locality for different set of tasks, by executing them on specific cores.

kThreads are kernel-level threads (std::thread), that are the main vehicle to utilize cores and execute the program. Each **kThread** can interact with the local scheduler in the **Cluster** and execute the *uThreads* provided by the local *Scheduler*, but it can move uThreads to another **Cluster** in the application. The former can happen when uThreads *yield* or *block* at user level, and the latter happens when uThreads *migrate* to another **Cluster**. Migration let the **uThread** continue execution on a different set of kThreads based on the requirements of the code.

uThreads are the main building blocks of the library. They are either sitting in a readyQueue or runQueue waiting to be picked by a **kThread**, running by a **kThread**, or blocked and waiting for an event to occur. uThreads are being

scheduled cooperatively over Clusters, they can either yield, migrate or block on an event to let other uThreads utilized the same kThread they are being executed over.

Each application has at least one Cluster, one kThread and one uThread. Each C++ application has at least one thread of execution (kernel thread) which runs the *main()* function. A C++ application that is linked with uThreads library, upon execution, creates a *defaultCluster*, a wrapper around the main execution thread and call it *defaultkThread*, and also a uThread called *mainUT* to take over the defaultkThread stack and run the *main* function.

In addition, there is a single Poller kThread which is responsible for polling the network devices, and multiplexing network events over the Cluster. uThreads provides a user-level blocking on network events, where network calls are non-blocking at the kernel-level but uThreads block on network events if the device is not ready. The poller thread is thus responsible for unblocking the uThreads upon receiving the related network event. The poller thread is using *edge triggered epoll* in Linux, and which is similar to how Golang supports multiplexing of network events.

Currently, uThreads only supports **fixed** stack sizes for performance purposes. uThread's stack is cached after finishing execution to avoid the extra overhead of memory allocation.

Scheduler

As Explained earlier, each Cluster has a local Scheduler which is responsible for distributing uThreads among the kThreads within that Cluster. Currently, there are 4 different schedulers, which will be explained below. These schedulers do not support any work sharing or work stealing at the moment and based on the type of the Scheduler either uThreads are assigned in a round robin manner, or kThreads ask the Scheduler for more uThreads when they run out of work. The type of the scheduler is determined at compile time by defining **SCHEDULERNO** and pass the related scheduler number. The default scheduler is Scheduler #2, with local intrusive Multiple-Producer-Single-Consumer Queues per kThread (since it provides better performance and scalability).

- **Global ReadyQueue per Cluster:** The first scheduler is implemented using an unbounded intrusive ReadyQueue per Cluster, and C++ synchronization primitives (`std::mutex` and `std::condition_variable`) are used to orchestrate kThreads to access the ReadyQueue. To avoid the high overhead of mutex and condition_variable (in linux: `pthread_mutex` and `pthread_cond`) under contention, a local queue is added to each kThread, so everytime a kThread runs out of work, it simply removes many uThreads, based on the size of the ReadyQueue and the number of kThreads in that Cluster, instead of one. The local queue is only accessed by a single kThread and thus does not require mutual exclusion. The following shows the design: To measure the performance of different schedulers, the following experiment is designed:
 - The experiment starts with 2 Clusters.
 - There are k kThreads per Cluster, and the number of kThreads are changed (x-axis of the following graphs shows the number of kThreads).
 - There are $1,000,000 \times k$ uThreads created when the experiment starts.
 - Each uThread migrates back and forth between Clusters.
 - uThread exits after 10 migrations.
 - We measure the number of migrations/second and plot it on y-axis, thus higher is better.

Figure 3 shows result for the first Scheduler:

Based on the results, this approach is not very scalable past 4 kThreads per Cluster.

- **Local RunQueue per kThread using mutex and cv:** To provide better scalability, we can remove the global ReadyQueue to avoid the contention for mutex. Thus, the next scheduler (#3, numbering does not follow the story line here), provides local unbounded intrusive queue per kThread and removes the global ReadyQueue. The scheduler assign the uThreads to kThreads in a round-robin manner. Each queue is protected with a `std::mutex` and `std::condition_variable`, and the following figure shows the design: and here are the results: Removing the bottleneck and getting rid of the central lock seems to provide better scalability, but can we do better?

- **Local RunQueue per kThread using lock-free non-intrusive Multiple-Producer-Single-Consumer Queue:** Since the only consumer for each local queue, is a single `kThread`, to reduce the synchronization overhead it is better to use a lock-free Multiple-Producer-Single-Consumer queue. The queue that is being used is a non-intrusive queue (you can find an implementation in the source code or [here](#)). With this queue, there is no contention on the consumer side and producers rarely block the consumer, and to push to the queue producers using an atomic exchange. Here is the result for using Scheduler #4:
- **Local RunQueue per kThread using lock-free intrusive Multiple-Producer-Single-Consumer Queue:** To avoid managing an extra state, the above queue is modified to be intrusive, and here are the results for using the intrusive queue:

For lower number of threads the global queue seems to do better, but as the number increases the intrusive lock-free MPSC Queue is a better Choice. The default scheduler for uThreads is Scheduler #2 (intrusive MPSCQ), which can be changed at compile time by defining **SCHEDULERNO** and set it to the appropriate scheduler.

You can add your own scheduler by looking at the source code under `src/runtime/schedulers`, and provide a scheduler number and a Scheduler class in its own header. Documentation for adding a new scheduler will be added as soon as the code base reaches a stable state.

Migration and Joinable uThreads

uThreads can be joinable, which means upon creation the creator `uThread` can block until they finish execution. There are three ways to execute a piece of code on another `Cluster` (These can be used to execute tasks asynchronously on current `Cluster` or a remote one):

- **Migration:** `uThread` can migrate to another `Cluster` to execute a piece of code and it can either migrate back to the previous `Cluster` or continue the execution on the same `Cluster` or migrate to a different `Cluster`. The state is saved in the stack and when migrated the state is resumed from the `uThread`'s stack. The following code demonstrates a simple scenario of migrating to a different cluster and back, assuming `uThread` is executing on the *defaultCluster*:

```
Cluster *cluster1;

void func(){
    // some code
    migrate(*cluster1);
    // code to run on cluster1
    migrate(Cluster::getDefaultCluster());
    // some more code
}

int main() {
    cluster1 = new Cluster();
    kThread kt(*cluster1);
    uThread *ut = uThread::create();
    ut->start(Cluster::getDefaultCluster(), func);
    .
    .
    .
}
```

- **Non-joinable uThread:** Create a non-joinable `uThread` on the remote `Cluster` and wait for it to finish execution. While waiting, the creator `uThread` continues execution and does not care about when the new `uThread` finishes the job.

```

Cluster *cluster1;

void run(){
    //code to run on cluster1
}
void func(){
    // some code
    uThread *ut2 = uThread::create(false); //create a non-joinable thread
    ut2->start(cluster1, run);
    //continue executing
    // some more code
}

int main(){

    cluster1 = new Cluster();
    kThread kt(*cluster1);
    uThread *ut = uThread::create();
    ut->start(Cluster::getDefaultCluster(), func);
    .
    .
    .
}

```

- **Joinable uThread:** Create a joinable thread on the remote [Cluster](#) and wait for it to finish execution. While waiting, the [uThread](#) is blocked at user-level and will be unblocked when the new [uThread](#) finishes its job. Currently, only the creator can wait on the new [uThread](#), waiting on the [uThread](#) from other uThreads leads to undefined behaviour. This will be fixed in the near future.

```

Cluster *cluster1;

void run(){
    //code to run on cluster1
}
void func(){
    // some code
    uThread *ut2 = uThread::create(true); //create a joinable thread
    ut2->start(cluster1, run);
    ut2->join(); //wait for ut2 to finish execution and join
    // some more code
}

int main(){

    cluster1 = new Cluster();
    kThread kt(*cluster1);
    uThread *ut = uThread::create();
    ut->start(Cluster::getDefaultCluster(), func);
    .
    .
    .
}

```

User-level Blocking Synchronization Primitives

uThreads also provides user-level blocking synchronization and mutex primitives. It has basic [Mutex](#), Condition Variable and [Semaphore](#). You can find examples of their usage under *test* directory.

Examples

You can find various examples under the test directory in the [github repo](#). There is an [EchoClient](#) and [EchoServer](#) implemented using uThreads.

There is also a simple [webserver](#) to test uThreads functionality.

For performance comparisons, memached code has been updated to use uThreads instead of event loops (except the thread that accepts connections), where tasks are assigned to uThreads instead of using the underlying event library. The code can be found [here](#).

Acknowledgement

This work made possible through invaluable helps and advices I received from [Martin Karsten](#).

Chapter 2

Hierarchical Index

2.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

BlockingQueue	13
Cluster	16
ConditionVariable	18
Connection	19
IOPoller	26
Link	
kThread	26
uThread	34
Mutex	29
OwnerLock	31
Semaphore	33
uThreadCache	39

Chapter 3

Class Index

3.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

BlockingQueue	A queue used to keep track of blocked uThreads	13
Cluster	Scheduler and Cluster of kThreads	16
ConditionVariable	A user level condition variable	18
Connection	Represents a network connection	19
IOPoller	26
kThread	Object to represent kernel threads	26
Mutex	A user-level Mutex	29
OwnerLock	Owner Mutex where owner can recursively acquire the Mutex	31
Semaphore	A user-level Semaphore	33
uThread	User-level threads (fiber)	34
uThreadCache	Data structure to cache uThreads	39

Chapter 4

File Index

4.1 File List

Here is a list of all documented files with brief descriptions:

/home/saman/Programming/Research/uThreads/src/ cwrapper.h	
C Wrapper for uThreads	41
/home/saman/Programming/Research/uThreads/src/io/ EpollIOHandler.h	??
/home/saman/Programming/Research/uThreads/src/io/ IOHandler.h	??
/home/saman/Programming/Research/uThreads/src/io/ Network.h	??
/home/saman/Programming/Research/uThreads/src/runtime/ BlockingSync.h	??
/home/saman/Programming/Research/uThreads/src/runtime/ Cluster.h	??
/home/saman/Programming/Research/uThreads/src/runtime/ kThread.h	??
/home/saman/Programming/Research/uThreads/src/runtime/ Stack.h	??
/home/saman/Programming/Research/uThreads/src/runtime/ uThread.h	??
/home/saman/Programming/Research/uThreads/src/runtime/ uThreadCache.h	??

Chapter 5

Class Documentation

5.1 BlockingQueue Class Reference

A queue used to keep track of blocked uThreads.

```
#include <BlockingSync.h>
```

Public Member Functions

- `BlockingQueue (const BlockingQueue &)=delete`
BlockingQueue cannot be copied or assigned.
- `const BlockingQueue & operator= (const BlockingQueue &)=delete`
BlockingQueue cannot be copied or assigned.
- `bool empty () const`
Whether the queue is empty or not.
- `bool suspend (std::mutex &lock)`
Suspends the uThread and add it to the queue.
- `bool suspend (Mutex &)`
Suspends the uThread and add it to the queue.
- `bool signal (std::mutex &lock, uThread *&owner)`
Unblock one blocked uThread, used for OwnerLock.
- `bool signal (std::mutex &lock)`
unblock one blocked, used for Mutex
- `bool signal (Mutex &)`
unblock one blocked uThread, used for ConditionVariable
- `void signalAll (Mutex &)`
unblock all blocked uThreads, used for Condition Variable

Static Public Member Functions

- `template<typename T >`
`static void postSwitchFunc (void *ut, void *args)`

Friends

- class **uThread**
- class **Mutex**
- class **OwnerLock**
- class **ConditionVariable**
- class **Semaphore**

5.1.1 Detailed Description

A queue used to keep track of blocked uThreads.

This queue is a FIFO queue used to hold blocked uThreads on [Mutex](#), [Semaphore](#), or Condition Variable.

5.1.2 Member Function Documentation

5.1.2.1 `bool BlockingQueue::signal (std::mutex & lock, uThread *& owner)`

Unblock one blocked [uThread](#), used for [OwnerLock](#).

Parameters

<i>lock</i>	mutex to be released after signal is done
<i>owner</i>	passed to support atomic setting of <code>Mutex::owner</code>

Returns

true if a [uThread](#) was unblocked, and false otherwise

5.1.2.2 `bool BlockingQueue::signal (std::mutex & lock) [inline]`

unblock one blocked, used for [Mutex](#)

Parameters

<i>lock</i>	mutex to be released after signal is done
-------------	---

Returns

true if a [uThread](#) was unblocked, and false otherwise

5.1.2.3 `bool BlockingQueue::signal (Mutex & mutex)`

unblock one blocked [uThread](#), used for [ConditionVariable](#)

Parameters

<i>Mutex</i>	that is released after signal is done
--------------	---------------------------------------

Returns

true if a [uThread](#) was unblocked, and false otherwise

5.1.2.4 void BlockingQueue::signalAll (*Mutex* & *mutex*)

unblock all blocked uThreads, used for Condition Variable

Parameters

<i>Mutex</i>	to be released after signalAll is done
--------------	--

5.1.2.5 bool BlockingQueue::suspend (std::mutex & *lock*)

Suspends the [uThread](#) and add it to the queue.

Parameters

<i>lock</i>	a mutex to be released after blocking
-------------	---------------------------------------

Returns

whether the suspension was successful or not

Suspends the [uThread](#) and adds it to the [BlockingQueue](#).

5.1.2.6 bool BlockingQueue::suspend (*Mutex* & *mutex*)

Suspends the [uThread](#) and add it to the queue.

Parameters

<i>lock</i>	a mutex to be released after blocking
-------------	---------------------------------------

Returns

whether the suspension was successful or not

Suspends the [uThread](#) and adds it to the [BlockingQueue](#).

The documentation for this class was generated from the following files:

- /home/saman/Programming/Research/uThreads/src/runtime/BlockingSync.h
- /home/saman/Programming/Research/uThreads/src/runtime/BlockingSync.cpp

5.2 Cluster Class Reference

Scheduler and [Cluster](#) of kThreads.

```
#include <Cluster.h>
```

Public Member Functions

- [Cluster](#) ()
- [Cluster](#) (const [Cluster](#) &)=delete
- const [Cluster](#) & [operator=](#) (const [Cluster](#) &)=delete
- uint64_t [getID](#) () const
Get the ID of [Cluster](#).
- size_t [getNumberOfkThreads](#) () const
Total number of kThreads belonging to this cluster.

Static Public Member Functions

- static [Cluster](#) & [getDefaultCluster](#) ()

Friends

- class **kThread**
- class **uThread**
- class **Connection**
- class **IOHandler**
- class **Scheduler**

5.2.1 Detailed Description

Scheduler and [Cluster](#) of kThreads.

[Cluster](#) is an entity that contains multiple kernel threads ([kThread](#)). Each cluster is responsible for maintaining a ready queue and performing basic scheduling tasks. Programs can have as many Clusters as is necessary. The [Cluster](#)'s ReadyQueue is a multiple-producer multiple-consumer queue where consumers are only kThreads belonging to that [Cluster](#), and producers can be any running [kThread](#). kThreads constantly push and pull uThreads to/from the ReadyQueue. [Cluster](#) is an interface between kThreads and the ReadyQueue, and also provides the means to group kThreads together.

Each [Cluster](#) has its own IOHandler. IOHandler is responsible for providing asynchronous nonblocking access to IO devices. For now each instance of an IOHandler has its own dedicated poller thread, which means each cluster has a dedicated IO poller thread when it is created. This might change in the future. Each [uThread](#) that requires access to IO uses the IOHandler to avoid blocking the [kThread](#), if the device is ready for read or write, the [uThread](#) continues otherwise it blocks until it is ready, and the [kThread](#) execute another [uThread](#) from the ReadyQueue.

When the program starts a defaultCluster is created for the kernel thread that runs the *main* function. defaultCluster can be used like any other clusters.

5.2.2 Constructor & Destructor Documentation

5.2.2.1 Cluster::Cluster ()

Create a new [Cluster](#)

5.2.2.2 Cluster::Cluster (const Cluster &) [delete]

[Cluster](#) cannot be copied or assigned.

5.2.3 Member Function Documentation

5.2.3.1 static Cluster& Cluster::getDefaultCluster () [inline],[static]

Returns

defaultCluster

5.2.3.2 uint64_t Cluster::getID () const [inline]

Get the ID of [Cluster](#).

Returns

The ID of the cluster

5.2.3.3 size_t Cluster::getNumberOfkThreads () const [inline]

Total number of kThreads belonging to this cluster.

Returns

Total number of kThreads belonging to this cluster

5.2.3.4 const Cluster& Cluster::operator= (const Cluster &) [delete]

[Cluster](#) cannot be copied or assigned.

The documentation for this class was generated from the following files:

- /home/saman/Programming/Research/uThreads/src/runtime/Cluster.h
- /home/saman/Programming/Research/uThreads/src/runtime/Cluster.cpp
- /home/saman/Programming/Research/uThreads/src/runtime/uThread.cpp

5.3 ConditionVariable Class Reference

A user level condition variable.

```
#include <BlockingSync.h>
```

Public Member Functions

- void `wait` (`Mutex` &mutex)
Block `uThread` on the condition variable using the provided mutex.
- void `signal` (`Mutex` &mutex)
Unblock one `uThread` waiting on the condition variable.
- void `signalAll` (`Mutex` &mutex)
unblock all `uThreads` waiting on the condition variable
- bool `empty` ()
Whether the waiting list is empty or not.

5.3.1 Detailed Description

A user level condition variable.

User-level Condition Variable blocks only in user-level by suspending the `uThreads` instead of blocking the kernel threads.

5.3.2 Member Function Documentation

5.3.2.1 bool ConditionVariable::empty () `[inline]`

Whether the waiting list is empty or not.

Returns

Whether the waiting list is empty or not

5.3.2.2 void ConditionVariable::signal (`Mutex` & *mutex*) `[inline]`

Unblock one `uThread` waiting on the condition variable.

Parameters

<i>mutex</i>	The mutex to be released after unblocking is done
--------------	---

5.3.2.3 void ConditionVariable::signalAll (`Mutex` & *mutex*) `[inline]`

unblock all `uThreads` waiting on the condition variable

Parameters

<i>mutex</i>	The mutex to be released after unblocking is done
--------------	---

5.3.2.4 void ConditionVariable::wait (Mutex & *mutex*) [inline]

Block [uThread](#) on the condition variable using the provided mutex.

Parameters

<i>mutex</i>	used to synchronize access to the condition
--------------	---

The documentation for this class was generated from the following file:

- /home/saman/Programming/Research/uThreads/src/runtime/BlockingSync.h

5.4 Connection Class Reference

Represents a network connection.

```
#include <Network.h>
```

Public Member Functions

- [Connection](#) ()
Create a [Connection](#) that does not have.
- [Connection](#) (int fd)
Create a connection object with the provided fd.
- [Connection](#) (int domain, int type, int protocol) throw (std::system_error)
Same as socket syscall adds | SOCK_NONBLOCK to type.
- int [accept](#) ([Connection](#) *conn, struct sockaddr *addr, socklen_t *addrlen)
nonblocking accept syscall and updating the passed [Connection](#) object
- [Connection](#) * [accept](#) (struct sockaddr *addr, socklen_t *addrlen) throw (std::system_error)
Accepts a connection and returns a connection object.
- int [socket](#) (int domain, int type, int protocol)
Same as socket syscall, set the fd for current connection.
- int [listen](#) (int backlog)
Same as listen syscall on current fd.
- int [bind](#) (const struct sockaddr *addr, socklen_t addrlen)
Same as bind syscall.
- int [connect](#) (const struct sockaddr *addr, socklen_t addrlen)
Same as connect syscall.
- ssize_t [recv](#) (void *buf, size_t len, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [recvfrom](#) (void *buf, size_t len, int flags, struct sockaddr *src_addr, socklen_t *addrlen)
Call the underlying system call on [Connection](#)'s file descriptor.

- `ssize_t recvmmsg` (int sockfd, struct msghdr *msg, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- `int recvmmsg` (int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags, struct timespec *timeout)
Call the underlying system call on [Connection](#)'s file descriptor.
- `ssize_t send` (const void *buf, size_t len, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- `ssize_t sendto` (int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen)
Call the underlying system call on [Connection](#)'s file descriptor.
- `ssize_t sendmsg` (const struct msghdr *msg, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- `int sendmmsg` (int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- `ssize_t read` (void *buf, size_t count)
Call the underlying system call on [Connection](#)'s file descriptor.
- `ssize_t write` (const void *buf, size_t count)
Call the underlying system call on [Connection](#)'s file descriptor.
- `void blockOnRead` ()
Block [uThread](#), waiting for fd to become ready for read.
- `void blockOnWrite` ()
Block [uThread](#), waiting for fd to become ready for write.
- `int close` ()
closes the socket
- `int getFd` () const
return the [Connection](#)'s file descriptor

5.4.1 Detailed Description

Represents a network connection.

[Connection](#) class is a wrapper around socket and provides the ability to do nonblocking read/write on sockets, and nonblocking accept. It first tries to read/write/accept and if the fd is not ready uses the underlying polling structure to wait for the fd to be ready. Thus, the [uThread](#) that is calling these functions is blocked if the fd is not ready, and [kThread](#) never blocks.

5.4.2 Constructor & Destructor Documentation

5.4.2.1 `Connection::Connection ()` `[inline]`

Create a [Connection](#) that does not have.

This is useful for accept or socket functions that require a [Connection](#) object without fd being set

5.4.2.2 `Connection::Connection (int fd)` `[inline]`

Create a connection object with the provided fd.

Parameters

<i>fd</i>	If the connection is already established by other means, set the fd and add it to the polling structure
-----------	---

5.4.2.3 `Connection::Connection (int domain, int type, int protocol) throw std::system_error`

Same as socket syscall adds | SOCK_NONBLOCK to type.

Returns

same as socket syscall

Throws a `std::system_error` exception. Do not call from C code. The underlying socket is always nonblocking. This is achieved by adding a (| SOCK_NONBLOCK) to type, thus requires linux kernels > 2.6.27

5.4.3 Member Function Documentation

5.4.3.1 `int Connection::accept (Connection * conn, struct sockaddr * addr, socklen_t * addrlen)`

nonblocking accept syscall and updating the passed [Connection](#) object

Parameters

<i>conn</i>	Pointer to a Connection object that is not initialized
-------------	--

Returns

same as accept system call

This format is used for compatibility with C

5.4.3.2 `Connection * Connection::accept (struct sockaddr * addr, socklen_t * addrlen) throw std::system_error`

Accepts a connection and returns a connection object.

Returns

Newly created connection

Throws a `std::system_error` exception on error. Never call from C.

5.4.3.3 `int Connection::bind (const struct sockaddr * addr, socklen_t addrlen)`

Same as bind syscall.

Returns

Same as bind syscall

5.4.3.4 `int Connection::close ()`

closes the socket

Returns

the same as close system call

5.4.3.5 `int Connection::connect (const struct sockaddr * addr, socklen_t addrlen)`

Same as connect syscall.

Returns

Same as connect syscall

5.4.3.6 `int Connection::getFd () const` `[inline]`

return the [Connection](#)'s file descriptor

Returns

file descriptor

5.4.3.7 `int Connection::listen (int backlog)`

Same as listen syscall on current fd.

Returns

Same as listen syscall

5.4.3.8 `ssize_t Connection::read (void * buf, size_t count)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.9 `ssize_t Connection::recv (void * buf, size_t len, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.10 `ssize_t Connection::recvfrom (void * buf, size_t len, int flags, struct sockaddr * src_addr, socklen_t * addrlen)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.11 `int Connection::recvmsg (int sockfd, struct mmsghdr * msgvec, unsigned int vlen, unsigned int flags, struct timespec * timeout)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.12 `ssize_t Connection::recvmsg (int sockfd, struct msghdr * msg, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.13 `ssize_t Connection::send (const void * buf, size_t len, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.14 `int Connection::sendmmsg (int sockfd, struct mmsghdr * msgvec, unsigned int vlen, unsigned int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.15 `ssize_t Connection::sendmsg (const struct msghdr * msg, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.16 `ssize_t Connection::sendto (int sockfd, const void * buf, size_t len, int flags, const struct sockaddr * dest_addr, socklen_t addrlen)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

5.4.3.17 `int Connection::socket (int domain, int type, int protocol)`

Same as socket syscall, set the fd for current connection.

Returns

same as socket syscall The underlying socket is always nonblocking. This is achieved by adding a `(| SOCK_↵ K_NONBLOCK)` to type, thus requires linux kernels > 2.6.27

5.4.3.18 ssize_t Connection::write (const void * buf, size_t count)

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: (res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

The documentation for this class was generated from the following files:

- /home/saman/Programming/Research/uThreads/src/io/Network.h
- /home/saman/Programming/Research/uThreads/src/io/Network.cpp

5.5 IOPoller Class Reference

Protected Member Functions

- **IOPoller** (IOHandler &)

The documentation for this class was generated from the following files:


- /home/saman/Programming/Research/uThreads/src/io/EpollIOHandler.h
- /home/saman/Programming/Research/uThreads/src/io/EpollIOHandler.cpp

5.6 kThread Class Reference


Object to represent kernel threads.

```
#include <kThread.h>
```

Inheritance diagram for kThread:



Collaboration diagram for kThread:



Public Member Functions

- `kThread (Cluster &)`
Create a `kThread` on the passed cluster.
- `kThread (const kThread &)=delete`
`kThread` cannot be copied or assigned.
- `const kThread & operator= (const kThread &)=delete`
`kThread` cannot be copied or assigned.
- `std::thread::native_handle_type getThreadNativeHandle ()`
return the native handle for the kernel thread
- `std::thread::id getID ()`
returns the kernel thread ID

Static Public Member Functions

- `static kThread * currentkThread ()`
Get the pointer to the current `kThread`.
- `static uint getTotalNumberOfThreads ()`

Friends

- class `uThread`
- class `Cluster`
- class `IOHandler`
- class `Scheduler`

5.6.1 Detailed Description

Object to represent kernel threads.

[kThread](#) is an interface for underlying kernel threads. kThreads are pulling and pushing uThreads from ReadyQueue provided by the [Cluster](#) and context switch to them and execute them. Each [kThread](#) belongs to only and only one [Cluster](#) and it can only pull uThreads from the ReadyQueue of that [Cluster](#). However, kThreads can push uThreads to the ReadyQueue of any cluster.

defaultKT is the first kernel thread that executes and is responsible for running the main() function. defaultKT is created when the program starts.

Each [kThread](#) has a mainUT which is a [uThread](#) used when the ReadyQueue is empty. mainUT is used to switch out the previous [uThread](#) and either pull uThreads from the ReadyQueue if it's not empty, or block on a condition variable waiting for uThreads to be pushed to the queue.

kThreads can be created by passing a [Cluster](#) to the constructor of the [kThread](#).

5.6.2 Constructor & Destructor Documentation

5.6.2.1 kThread::kThread ([Cluster](#) & *cluster*)

Create a [kThread](#) on the passed cluster.

Parameters

The	Cluster this kThread belongs to.
-----	--

5.6.3 Member Function Documentation

5.6.3.1 static [kThread](#)* [kThread::currentkThread](#) () `[inline], [static]`

Get the pointer to the current [kThread](#).

Returns

current [kThread](#)

This is necessary when a [uThread](#) wants to find which [kThread](#) it is being executed over.

5.6.3.2 `std::thread::id` [kThread::getID](#) ()

returns the kernel thread ID

Returns

the [kThread](#) ID

The returned type depends on the platform.

5.6.3.3 `std::thread::native_handle_type kThread::getThreadNativeHandle ()`

return the native handle for the kernel thread

Returns

native handle for the [kThread](#)

In linux this is `pthread_t` representation of the thread.

5.6.3.4 `static uint kThread::getTotalNumberOfkThreads () [inline],[static]`

Returns

total number of `kThreads` running under the program.

5.6.3.5 `const kThread& kThread::operator= (const kThread &) [delete]`

[kThread](#) cannot be copied or assigned.

The documentation for this class was generated from the following files:


- `/home/saman/Programming/Research/uThreads/src/runtime/kThread.h`
- `/home/saman/Programming/Research/uThreads/src/runtime/kThread.cpp`
- `/home/saman/Programming/Research/uThreads/src/runtime/uThread.cpp`

5.7 Mutex Class Reference

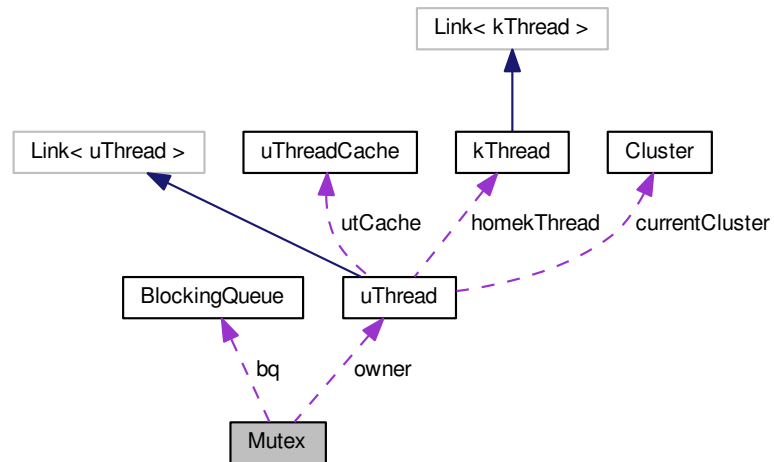
A user-level [Mutex](#).

```
#include <BlockingSync.h>
```

Inheritance diagram for `Mutex`:



Collaboration diagram for Mutex:



Public Member Functions

- template<bool OL = false>
bool **acquire** ()
acquire the mutex
- void **release** ()
release the `Mutex`
- void **unlock** ()

Protected Member Functions

- template<bool OL>
bool **internalAcquire** (mword timeout)
- void **internalRelease** ()

Protected Attributes

- std::mutex **lock**
- `BlockingQueue` **bq**
- `uThread` * **owner**

Friends

- class **ConditionVariable**
- class **Semaphore**
- class **BlockingQueue**

5.7.1 Detailed Description

A user-level [Mutex](#).

5.7.2 Member Function Documentation

5.7.2.1 `template<bool OL = false> bool Mutex::acquire () [inline]`

acquire the mutex

Returns

true if it was acquired, false otherwise

The return value is only for when timeouts are implemented

The documentation for this class was generated from the following file:


- `/home/saman/Programming/Research/uThreads/src/runtime/BlockingSync.h`

5.8 OwnerLock Class Reference

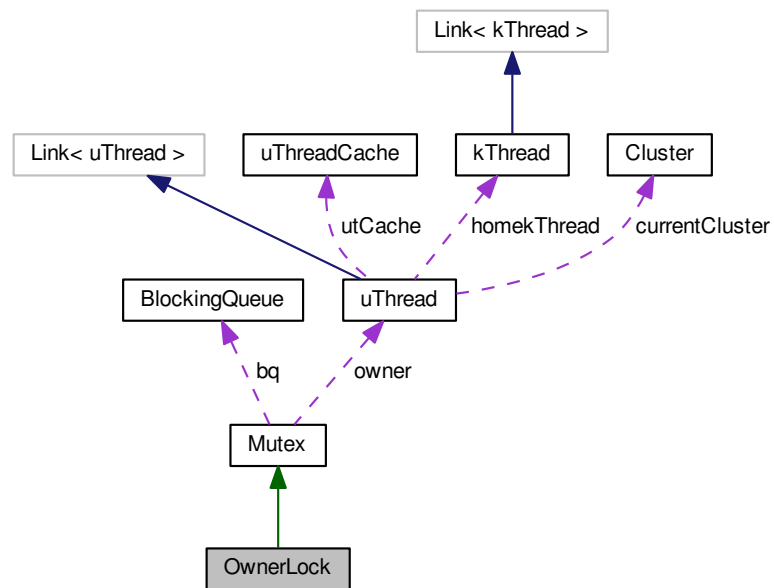
an Owner [Mutex](#) where owner can recursively acquire the [Mutex](#)

```
#include <BlockingSync.h>
```

Inheritance diagram for OwnerLock:



Collaboration diagram for OwnerLock:



Public Member Functions

- `OwnerLock ()`
Create a new *OwnerLock*.
- `mword acquire ()`
Acquire the *OwnerLock*.
- `mword release ()`
Release the *OwnerLock* once.

Additional Inherited Members

5.8.1 Detailed Description

an Owner `Mutex` where owner can recursively acquire the `Mutex`

5.8.2 Member Function Documentation

5.8.2.1 `mword OwnerLock::acquire () [inline]`

Acquire the *OwnerLock*.

Returns

The number of times current owner acquired the lock

5.8.2.2 mword OwnerLock::release () [inline]

Release the [OwnerLock](#) once.

Returns

The number of times current owner acquired the lock, if lock is released completely the result is 0

The documentation for this class was generated from the following file:

- /home/saman/Programming/Research/uThreads/src/runtime/BlockingSync.h

5.9 Semaphore Class Reference

A user-level [Semaphore](#).

```
#include <BlockingSync.h>
```

Public Member Functions

- [Semaphore](#) (mword c=0)
Create a new [Semaphore](#).
- bool [P](#) ()
Decrement the value of the [Semaphore](#).
- void [V](#) ()
increment the value of the [Semaphore](#)

5.9.1 Detailed Description

A user-level [Semaphore](#).

blocks in user level by blocking uThreads and does not cause kernel threads to block.

5.9.2 Constructor & Destructor Documentation

5.9.2.1 Semaphore::Semaphore (mword c = 0) [inline], [explicit]

Create a new [Semaphore](#).

Parameters

<i>c</i>	Initial value of the Semaphore
----------	--

5.9.3 Member Function Documentation

5.9.3.1 `bool Semaphore::P () [inline]`

Decrement the value of the [Semaphore](#).

Returns

Whether it was successful or not

The documentation for this class was generated from the following file:


- `/home/saman/Programming/Research/uThreads/src/runtime/BlockingSync.h`

5.10 uThread Class Reference

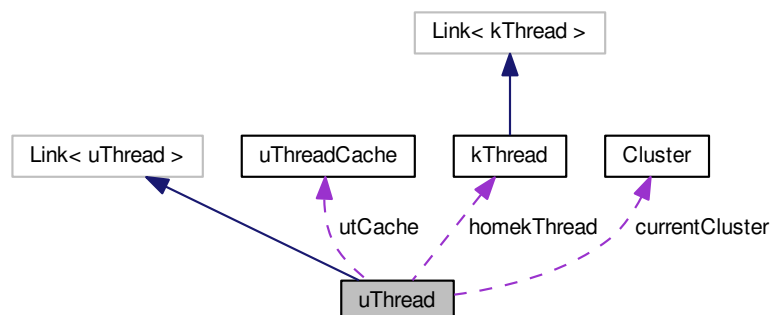
user-level threads (fiber)

```
#include <uThread.h>
```

Inheritance diagram for uThread:



Collaboration diagram for uThread:



Public Member Functions

- `uThread` (const `uThread` &)=delete
`uThread` cannot be copied or assigned
- const `uThread` & `operator=` (const `uThread` &)=delete
- void `start` (const `Cluster` &cluster, ptr_t func, ptr_t arg1=nullptr, ptr_t arg2=nullptr, ptr_t arg3=nullptr)
start the `uThread` by calling the function passed to it
- void `resume` ()
Resumes the `uThread`. If `uThread` is blocked or is waiting on IO it will be placed back on the ReadyQueue.
- bool `join` ()
Wait for `uThread` to finish execution and exit.
- void `detach` ()
Detach a joinable thread.
- `Cluster` & `getCurrentCluster` () const
return the current `Cluster` `uThread` is executed on
- uint64_t `getID` () const
get the ID of this `uThread`

Static Public Member Functions

- static `uThread` * `create` (size_t ss, bool joinable=false)
Create a `uThread` with a given stack size.
- static `uThread` * `create` (bool joinable=false)
Create a `uThread` with default stack size.
- static void `yield` ()
Causes `uThread` to yield.
- static void `terminate` ()
Terminates the `uThread`.
- static void `migrate` (`Cluster` *)
Move the `uThread` to the provided cluster.
- static uint64_t `getTotalNumberofUTs` ()
- static `uThread` * `currentUThread` ()
Get a pointer to the current running `uThread`.

Protected Types

- enum **State** : std::uint8_t {
 INITIALIZED, **READY**, **RUNNING**, **YIELD**,
 MIGRATE, **WAITING**, **TERMINATED** }
- enum **JState** : std::uint8_t { **DETACHED**, **JOINABLE**, **JOINING** }

Protected Member Functions

- `uThread` (vaddr sb, size_t ss)
- virtual void **destory** (bool)
- void **reset** ()
- void **suspend** (funcvoid2_t, void *)

Static Protected Member Functions

- static vaddr **createStack** (size_t)
- static void **invoke** (funcvoid3_t, ptr_t, ptr_t, ptr_t) __noreturn

Protected Attributes

- [Cluster](#) * **currentCluster**
- [kThread](#) * **homekThread**
- vaddr **stackPointer**
- vaddr **stackBottom**
- size_t **stackSize**
- UTVar * **utvar**
- uint64_t **uThreadID**
- enum uThread::State **state**
- enum uThread::JState **jState**

Static Protected Attributes

- static [uThreadCache](#) **utCache**
- static std::atomic_ulong **totalNumberOfUTs**
- static std::atomic_ulong **uThreadMasterID**

Friends

- class **uThreadCache**
- class **kThread**
- class **Cluster**
- class **BlockingQueue**
- class **IOHandler**
- class **Scheduler**

5.10.1 Detailed Description

user-level threads (fiber)

uThreads are building blocks of this library. They are lightweight threads and do not have the same context switch overhead as kernel threads. Each [uThread](#) is an execution unit provided to run small tasks. uThreads are being managed cooperatively and there is no preemption involved. uThreads either yield, migrate, or blocked and giving way to other uThreads to get a chance to run.

Due to the cooperative nature of uThreads, it is recommended that uThreads do not block the underlying kernel thread for a long time. However, since there can be multiple kernel threads ([kThread](#)) in the program, if one or more uThreads block underlying kThreads for small amount of time the execution of the program does not stop and other kThreads keep executing.

Another pitfall can be when all uThreads are blocked and each waiting for an event to occurs which can cause deadlocks. It is programmer's responsibility to make sure this never happens. Although it never happens unless a [uThread](#) is blocked without a reason (not waiting on a lock or IO), otherwise there is always an event (another [uThread](#) or polling structure) that unblock the [uThread](#).

Each [uThread](#) has its own stack which is very smaller than kernel thread's stack. This stack is allocated when [uThread](#) is created.

5.10.2 Constructor & Destructor Documentation

5.10.2.1 uThread::uThread (vaddr *sb*, size_t *ss*) [inline],[protected]

The main and only constructor for [uThread](#). uThreads are not supposed to be created by using the constructor. The memory used to save the [uThread](#) object is allocated at the beginning of its own stack. Thus, by freeing the stack memory [uThread](#) object is being destroyed as well. Therefore, the implicit destructor is not necessary.

5.10.3 Member Function Documentation

5.10.3.1 uThread * uThread::create (size_t *ss*, bool *joinable* = false) [static]

Create a [uThread](#) with a given stack size.

Parameters

<i>ss</i>	stack size
<i>joinable</i>	Whether this thread is joinable or detached

Returns

a pointer to a new [uThread](#)

This function relies on a [uThreadCache](#) structure and does not always allocate the stack.

5.10.3.2 static uThread* uThread::create (bool *joinable* = false) [inline],[static]

Create a [uThread](#) with default stack size.

Parameters

<i>joinable</i>	Whether this thread is joinable or detached
-----------------	---

Returns

a pointer to a new [uThread](#)

5.10.3.3 uThread * uThread::currentUThread () [static]

Get a pointer to the current running [uThread](#).

Returns

pointer to the current [uThread](#)

5.10.3.4 `Cluster& uThread::getCurrentCluster () const` `[inline]`

return the current `Cluster uThread` is executed on

Returns

the current `Cluster uThread` is executed on

5.10.3.5 `uint64_t uThread::getID () const` `[inline]`

get the ID of this `uThread`

Returns

ID of the `uThread`

5.10.3.6 `static uint64_t uThread::getTotalNumberOfUTs ()` `[inline],[static]`

Returns

Total number of `uThreads` in the program

This number does not include `mainUT` or `IOUTs`

5.10.3.7 `bool uThread::join ()`

Wait for `uThread` to finish execution and exit.

Returns

Whether join was successful or failed

5.10.3.8 `void uThread::migrate (Cluster * cluster)` `[static]`

Move the `uThread` to the provided cluster.

Parameters

<i>cluster</i>	This function is used to migrate the <code>uThread</code> to another <code>Cluster</code> . Migration is useful specially if clusters form a pipeline of execution.
----------------	---

5.10.3.9 `const uThread& uThread::operator= (const uThread &)` `[delete]`

5.10.3.10 `void uThread::start (const Cluster & cluster, ptr_t func, ptr_t arg1 = nullptr, ptr_t arg2 = nullptr, ptr_t arg3 = nullptr)`

start the `uThread` by calling the function passed to it

Parameters

<i>cluster</i>	The cluster that function belongs to.
<i>func</i>	a pointer to a function that should be executed by the <code>uThread</code> .
<i>arg1</i>	first argument of the function (can be <code>nullptr</code>)
<i>arg2</i>	second argument of the function (can be <code>nullptr</code>)
<i>arg3</i>	third argument of the function (can be <code>nullptr</code>)

After creating the `uThread` and allocating the stack, the `start()` function should be called to get the `uThread` going.

5.10.3.11 `void uThread::terminate () [static]`

Terminates the `uThread`.

By calling this function `uThread` is being terminated and `uThread` object is either destroyed or put back into the cache.

5.10.3.12 `void uThread::yield () [static]`

Causes `uThread` to yield.

`uThread` give up the execution context and place itself back on the ReadyQueue of the `Cluster`. If there is no other `uThreads` available to switch to, the current `uThread` continues execution.

The documentation for this class was generated from the following files:

- `/home/saman/Programming/Research/uThreads/src/runtime/uThread.h`
- `/home/saman/Programming/Research/uThreads/src/runtime/uThread.cpp`

5.11 uThreadCache Class Reference

Data structure to cache `uThreads`.

```
#include <uThreadCache.h>
```

Public Member Functions

- `uThreadCache` (size_t size=defaultuThreadCacheSize)
- `ssize_t push (uThread *ut)`
adds a `uThread` to the cache
- `uThread * pop ()`
pop a `uThread` from the list in FIFO order and return it

5.11.1 Detailed Description

Data structure to cache uThreads.

[uThreadCache](#) is a linked list of uThreads using and intrusive container to cache all terminated uThreads. Instead of destroying the memory allocated for the stack, simply reset the stack pointer and push it to the cache.

5.11.2 Member Function Documentation

5.11.2.1 `uThread* uThreadCache::pop ()` [inline]

pop a [uThread](#) from the list in FIFO order and return it

Returns

nullptr on failure, or a pointer to a [uThread](#) on success

5.11.2.2 `ssize_t uThreadCache::push (uThread * ut)` [inline]

adds a [uThread](#) to the cache

Parameters

<i>ut</i>	pointer to a uThread
-----------	--------------------------------------

Returns

size of the cache if push was successful or -1 if not

This function tries to push a [uThread](#) into the cache structure. If the cache is full or the mutex cannot be acquired immediately the operation has failed and the function returns -1. Otherwise, it adds the [uThread](#) to the list and return the size of the cache.

The documentation for this class was generated from the following file:

- `/home/saman/Programming/Research/uThreads/src/runtime/uThreadCache.h`


Chapter 6

File Documentation

6.1 /home/saman/Programming/Research/uThreads/src/cwrapper.h File Reference

C Wrapper for uThreads.

```
#include <pthread.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <inttypes.h>
#include <stdbool.h>
Include dependency graph for cwrapper.h:
```



Cluster

C interface for class [Cluster](#).

- typedef struct WCluster **WCluster**
- WCluster * [cluster_create](#) ()
- void [cluster_destroy](#) (WCluster *cluster)
- WCluster * [cluster_get_default](#) ()
- WCluster * [cluster_get_current](#) ()
- uint64_t [cluster_get_id](#) (WCluster *cluster)
Get the ID of [Cluster](#).
- size_t [cluster_get_number_of_kThreads](#) (WCluster *cluster)
Total number of kThreads belonging to this cluster.

kThread

C interface for class [kThread](#).

- typedef struct WkThread **WkThread**
- WkThread * [kThread_create](#) (WCluster *cluster)
- void [kThread_destroy](#) (WkThread *kt)
- uint64_t [kThread_get_total_number_of_kThreads](#) ()
- WkThread * [kThread_get_current](#) ()
Get the pointer to the current [kThread](#).
- pthread_t [kThread_get_current_pthread_id](#) ()
return the native handle for the kernel thread
- pthread_t [kThread_get_pthread_id](#) (WkThread *kt)
return the native handle for the kernel thread

uThread

C interface for class [uThread](#).

- typedef struct WuThread **WuThread**
- WuThread * [uThread_create](#) (bool joinable)
Create a [uThread](#) with a given stack size.
- void [uThread_start](#) (WuThread *ut, WCluster *cluster, void *func, void *arg1, void *arg2, void *arg3)
start the [uThread](#) by calling the function passed to it
- void [uThread_migrate](#) (WCluster *cluster)
Move the [uThread](#) to the provided cluster.
- void [uThread_terminate](#) (WuThread *ut)
Terminates the [uThread](#).
- void [uThread_yield](#) ()
Causes [uThread](#) to yield.
- bool [uThread_join](#) (WuThread *ut)
Wait for [uThread](#) to finish execution and exit.
- void [uThread_detach](#) (WuThread *ut)
Detach a joinable thread.
- uint64_t [uThread_get_id](#) (WuThread *ut)
get the ID of this [uThread](#)
- WuThread * [uThread_get_current](#) ()
Get a pointer to the current running [uThread](#).
- uint64_t [uThread_get_total_number_of_uThreads](#) ()

Connection

C interface for class [Connection](#).

- typedef struct WConnection **WConnection**
- WConnection * [connection_create](#) ()
Create a [Connection](#) that does not have.
- WConnection * [connection_create_with_fd](#) (int fd)
Create a connection object with the provided fd.
- WConnection * [connection_create_socket](#) (int domain, int type, int protocol)
Same as socket syscall adds | SOCK_NONBLOCK to type.
- void [connection_destroy](#) (WConnection *c)
- int [connection_accept](#) (WConnection *acceptor, WConnection *conn, struct sockaddr *addr, socklen_t *addrlen)
nonblocking accept syscall and updating the passed [Connection](#) object
- WConnection * [connection_accept_connection](#) (WConnection *acceptor, struct sockaddr *addr, socklen_t *addrlen)
Accepts a connection and returns a connection object.
- int [connection_socket](#) (WConnection *conn, int domain, int type, int protocol)
Same as socket syscall, set the fd for current connection.
- int [connection_listen](#) (WConnection *conn, int backlog)
Same as listen syscall on current fd.
- int [connection_bind](#) (WConnection *conn, const struct sockaddr *addr, socklen_t addrlen)
Same as bind syscall.
- int [connection_connect](#) (WConnection *conn, const struct sockaddr *addr, socklen_t addrlen)
Same as connect syscall.
- int [connection_close](#) (WConnection *conn)
closes the socket
- ssize_t [connection_recv](#) (WConnection *conn, void *buf, size_t len, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [connection_recvfrom](#) (WConnection *conn, void *buf, size_t len, int flags, struct sockaddr *src_addr, socklen_t *addrlen)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [connection_recvmsg](#) (WConnection *conn, int sockfd, struct msghdr *msg, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- int [connection_recvmsgvec](#) (WConnection *conn, int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags, struct timespec *timeout)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [connection_send](#) (WConnection *conn, const void *buf, size_t len, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [connection_sendto](#) (WConnection *conn, int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [connection_sendmsg](#) (WConnection *conn, const struct msghdr *msg, int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- int [connection_sendmsgvec](#) (WConnection *conn, int sockfd, struct mmsghdr *msgvec, unsigned int vlen, unsigned int flags)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [connection_read](#) (WConnection *conn, void *buf, size_t count)
Call the underlying system call on [Connection](#)'s file descriptor.
- ssize_t [connection_write](#) (WConnection *conn, const void *buf, size_t count)

Call the underlying system call on [Connection](#)'s file descriptor.

- void [connection_block_on_read](#) (WConnection *conn)
Block [uThread](#), waiting for fd to become ready for read.
- void [connection_block_on_write](#) (WConnection *conn)
Block [uThread](#), waiting for fd to become ready for write.
- int [connection_get_fd](#) (WConnection *conn)

Mutex

C interface for class [Mutex](#).

- typedef struct WMutex **WMutex**
- WMutex * [mutex_create](#) ()
- void [mutex_destroy](#) (WMutex *mutex)
- bool [mutex_acquire](#) (WMutex *mutex)
acquire the mutex
- void [mutex_release](#) (WMutex *mutex)
release the [Mutex](#)

OwnerLock

C interface for class [OwnerLock](#).

- typedef struct WOwnerLock **WOwnerLock**
- WOwnerLock * [ownerlock_create](#) ()
- void [ownerlock_destroy](#) (WOwnerLock *olock)
- uint64_t [ownerlock_acquire](#) (WOwnerLock *olock)
- void [ownerlock_release](#) (WOwnerLock *olock)

ConditionVariable

C interface for class [ConditionVariable](#).

- typedef struct WConditionVariable **WConditionVariable**
- WConditionVariable * [condition_variable_create](#) ()
- void [condition_variable_destroy](#) (WConditionVariable *cv)
- void [condition_variable_wait](#) (WConditionVariable *cv, WMutex *mutex)
Block [uThread](#) on the condition variable using the provided mutex.
- void [condition_variable_signal](#) (WConditionVariable *cv, WMutex *mutex)
Unblock one [uThread](#) waiting on the condition variable.
- void [condition_variable_signal_all](#) (WConditionVariable *cv, WMutex *mutex)
unblock all [uThreads](#) waiting on the condition variable
- bool [condition_variable_empty](#) (WConditionVariable *cv)
Whether the waiting list is empty or not.

Semaphore

C interface for class [Semaphore](#).

- typedef struct WSemaphore **WSemaphore**
- WSemaphore * [semaphore_create](#) ()
Create a new [Semaphore](#).
- void [semaphore_destroy](#) (WSemaphore *sem)
- bool [semaphore_p](#) (WSemaphore *sem)
Decrement the value of the [Semaphore](#).
- void [semaphore_v](#) (WSemaphore *sem)
increment the value of the [Semaphore](#)

uThreadPool

C interface for class uThreadPool.

- typedef struct WuThreadPool **WuThreadPool**
- WuThreadPool * [uthreadpool_create](#) ()
- void [uthreadpool_destory](#) (WuThreadPool *utp)
- void [uthreadpool_execute](#) (WuThreadPool *utp, WCluster *cluster, void *(*start_routine)(void *), void *arg)

6.1.1 Detailed Description

C Wrapper for uThreads.

Author

Saman Barghi

6.1.2 Function Documentation

6.1.2.1 WCluster* cluster_create ()

Create a new [Cluster](#)

6.1.2.2 void cluster_destroy (WCluster * cluster)

6.1.2.3 WCluster* cluster_get_current ()

6.1.2.4 WCluster* cluster_get_default ()

Returns

defaultCluster

6.1.2.5 uint64_t cluster_get_id (WCluster * *cluster*)

Get the ID of [Cluster](#).

Returns

The ID of the cluster

6.1.2.6 size_t cluster_get_number_of_kThreads (WCluster * *cluster*)

Total number of kThreads belonging to this cluster.

Returns

Total number of kThreads belonging to this cluster

6.1.2.7 WConditionVariable* condition_variable_create ()

6.1.2.8 void condition_variable_destroy (WConditionVariable * *cv*)

6.1.2.9 bool condition_variable_empty (WConditionVariable * *cv*)

Whether the waiting list is empty or not.

Returns

Whether the waiting list is empty or not

6.1.2.10 void condition_variable_signal (WConditionVariable * *cv*, WMutex * *mutex*)

Unblock one [uThread](#) waiting on the condition variable.

Parameters

<i>mutex</i>	The mutex to be released after unblocking is done
--------------	---

6.1.2.11 void condition_variable_signal_all (WConditionVariable * *cv*, WMutex * *mutex*)

unblock all uThreads waiting on the condition variable

Parameters

<i>mutex</i>	The mutex to be released after unblocking is done
--------------	---

6.1.2.12 `void condition_variable_wait (WConditionVariable * cv, WMutex * mutex)`

Block [uThread](#) on the condition variable using the provided mutex.

Parameters

<i>mutex</i>	used to synchronize access to the condition
--------------	---

6.1.2.13 `int connection_accept (WConnection * acceptor, WConnection * conn, struct sockaddr * addr, socklen_t * addrlen)`

nonblocking accept syscall and updating the passed [Connection](#) object

Parameters

<i>conn</i>	Pointer to a Connection object that is not initialized
-------------	--

Returns

same as accept system call

This format is used for compatibility with C

6.1.2.14 `WConnection* connection_accept_connexion (WConnection * acceptor, struct sockaddr * addr, socklen_t * addrlen)`

Accepts a connection and returns a connection object.

Returns

Newly created connection

Throws a `std::system_error` exception on error. Never call from C.

6.1.2.15 `int connection_bind (WConnection * conn, const struct sockaddr * addr, socklen_t addrlen)`

Same as bind syscall.

Returns

Same as bind syscall

6.1.2.16 `void connection_block_on_read (WConnection * conn)`

Block [uThread](#), waiting for fd to become ready for read.

6.1.2.17 `void connection_block_on_write (WConnection * conn)`

Block [uThread](#), waiting for fd to become ready for write.

6.1.2.18 `int connection_close (WConnection * conn)`

closes the socket

Returns

the same as close system call

6.1.2.19 `int connection_connect (WConnection * conn, const struct sockaddr * addr, socklen_t addrlen)`

Same as connect syscall.

Returns

Same as connect syscall

6.1.2.20 `WConnection* connection_create ()`

Create a [Connection](#) that does not have.

This is useful for accept or socket functions that require a [Connection](#) object without fd being set

6.1.2.21 `WConnection* connection_create_socket (int domain, int type, int protocol)`

Same as socket syscall adds | SOCK_NONBLOCK to type.

Returns

same as socket syscall

Throws a `std::system_error` exception. Do not call from C code. The unerlying socket is always nonblocking. This is achieved by adding a (`| SOCK_NONBLOCK`) to type, thus requires linux kernels > 2.6.27

6.1.2.22 `WConnection* connection_create_with_fd (int fd)`

Create a connection object with the provided fd.

Parameters

<i>fd</i>	If the connection is already established by other means, set the fd and add it to the polling structure
-----------	---

6.1.2.23 `void connection_destroy (WConnection * c)`

6.1.2.24 `int connection_get_fd (WConnection * conn)`

6.1.2.25 `int connection_listen (WConnection * conn, int backlog)`

Same as listen syscall on current fd.

Returns

Same as listen syscall

6.1.2.26 `ssize_t connection_read (WConnection * conn, void * buf, size_t count)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.27 `ssize_t connection_recv (WConnection * conn, void * buf, size_t len, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.28 `ssize_t connection_recvfrom (WConnection * conn, void * buf, size_t len, int flags, struct sockaddr * src_addr, socklen_t * addrlen)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related systemcall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.29 `int connection_recvmsg (WConnection * conn, int sockfd, struct mmsghdr * msgvec, unsigned int vlen, unsigned int flags, struct timespec * timeout)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related systemcall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.30 `ssize_t connection_recvmsg (WConnection * conn, int sockfd, struct msghdr * msg, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related systemcall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.31 `ssize_t connection_send (WConnection * conn, const void * buf, size_t len, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.32 `int connection_sendmmsg (WConnection * conn, int sockfd, struct mmsghdr * msgvec, unsigned int vlen, unsigned int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.33 `ssize_t connection_sendmsg (WConnection * conn, const struct msghdr * msg, int flags)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.34 `ssize_t connection_sendto (WConnection * conn, int sockfd, const void * buf, size_t len, int flags, const struct sockaddr * dest_addr, socklen_t addrlen)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.35 `int connection_socket (WConnection * conn, int domain, int type, int protocol)`

Same as socket syscall, set the fd for current connection.

Returns

same as socket syscall The underlying socket is always nonblocking. This is achieved by adding a `(| SOCK_↔ K_NONBLOCK)` to type, thus requires linux kernels > 2.6.27

6.1.2.36 `ssize_t connection_write (WConnection * conn, const void * buf, size_t count)`

Call the underlying system call on [Connection](#)'s file descriptor.

Returns

Same as what the related syscall returns

This function calls the system call with the same name. If the socket is ready for the required function it returns immediately, otherwise it blocks in the user-level (blocks [uThread](#) not [kThread](#)), and polls the file descriptor until it becomes ready.

The return results is the same as the underlying system call except that the following condition is never true when the function returns: `(res == -1) && (errno == EAGAIN || errno == EWOULDBLOCK)`

which means the [Connection](#) object does the polling and only returns when an error occurs or the socket is ready.

6.1.2.37 `WkThread* kThread_create (WCluster * cluster)`

6.1.2.38 `void kThread_destroy (WkThread * kt)`

6.1.2.39 `WkThread* kThread_get_current ()`

Get the pointer to the current [kThread](#).

Returns

current [kThread](#)

This is necessary when a [uThread](#) wants to find which [kThread](#) it is being executed over.

6.1.2.40 `pthread_t kThread_get_current_pthread_id ()`

return the native handle for the kernel thread

Returns

native handle for the [kThread](#)

In linux this is `pthread_t` representation of the thread.

6.1.2.41 `pthread_t kThread_get_pthread_id (WkThread * kt)`

return the native handle for the kernel thread

Returns

native handle for the [kThread](#)

In linux this is `pthread_t` representation of the thread.

6.1.2.42 `uint64_t kThread_get_total_number_of_kThreads ()`

Returns

total number of `kThreads` running under the program.

6.1.2.43 `bool mutex_acquire (WMutex * mutex)`

acquire the mutex

Returns

true if it was acquired, false otherwise

The return value is only for when timeouts are implemented

6.1.2.44 WMutex* mutex_create ()

6.1.2.45 void mutex_destroy (WMutex * mutex)

6.1.2.46 void mutex_release (WMutex * mutex)

release the [Mutex](#)

6.1.2.47 uint64_t ownerlock_acquire (WOwnerLock * olock)

6.1.2.48 WOwnerLock* ownerlock_create ()

6.1.2.49 void ownerlock_destroy (WOwnerLock * olock)

6.1.2.50 void ownerlock_release (WOwnerLock * olock)

6.1.2.51 WSemaphore* semaphore_create ()

Create a new [Semaphore](#).

Parameters

c	Initial value of the Semaphore
---	--

6.1.2.52 void semaphore_destroy (WSemaphore * sem)

6.1.2.53 bool semaphore_p (WSemaphore * sem)

Decrement the value of the [Semaphore](#).

Returns

Whether it was successful or not

6.1.2.54 void semaphore_v (WSemaphore * sem)

increment the value of the [Semaphore](#)

6.1.2.55 WuThread* uThread_create (bool joinable)

Create a [uThread](#) with a given stack size.

Parameters

<i>ss</i>	stack size
<i>joinable</i>	Whether this thread is joinable or detached

Returns

a pointer to a new [uThread](#)

This function relies on a [uThreadCache](#) structure and does not always allocate the stack.

6.1.2.56 `void uThread_detach (WuThread * ut)`

Detach a joinable thread.

6.1.2.57 `WuThread* uThread_get_current ()`

Get a pointer to the current running [uThread](#).

Returns

pointer to the current [uThread](#)

6.1.2.58 `uint64_t uThread_get_id (WuThread * ut)`

get the ID of this [uThread](#)

Returns

ID of the [uThread](#)

6.1.2.59 `uint64_t uThread_get_total_number_of_uThreads ()`

Returns

Total number of uThreads in the program

This number does not include mainUT or IOUts

6.1.2.60 `bool uThread_join (WuThread * ut)`

Wait for [uThread](#) to finish execution and exit.

Returns

Whether join was successful or failed

6.1.2.61 `void uThread_migrate (WCluster * cluster)`

Move the [uThread](#) to the provided cluster.

Parameters

<i>cluster</i>	This function is used to migrate the uThread to another Cluster . Migration is useful specially if clusters form a pipeline of execution.
----------------	---

6.1.2.62 `void uThread_start (WuThread * ut, WCluster * cluster, void * func, void * arg1, void * arg2, void * arg3)`

start the [uThread](#) by calling the function passed to it

Parameters

<i>cluster</i>	The cluster that function belongs to.
<i>func</i>	a pointer to a function that should be executed by the uThread .
<i>arg1</i>	first argument of the function (can be nullptr)
<i>arg2</i>	second argument of the function (can be nullptr)
<i>arg3</i>	third argument of the function (can be nullptr)

After creating the [uThread](#) and allocating the stack, the start() function should be called to get the [uThread](#) going.

6.1.2.63 `void uThread_terminate (WuThread * ut)`

Terminates the [uThread](#).

By calling this function [uThread](#) is being terminated and [uThread](#) object is either destroyed or put back into the cache.

6.1.2.64 `void uThread_yield ()`

Causes [uThread](#) to yield.

[uThread](#) give up the execution context and place itself back on the ReadyQueue of the [Cluster](#). If there is no other [uThreads](#) available to switch to, the current [uThread](#) continues execution.

Index

/home/saman/Programming/Research/uThreads/src/cwrapper.h, [41](#)

accept
 Connection, [21](#)

acquire
 Mutex, [31](#)
 OwnerLock, [32](#)

bind
 Connection, [21](#)

BlockingQueue, [13](#)
 signal, [14](#)
 signalAll, [15](#)
 suspend, [15](#)

close
 Connection, [21](#)

Cluster, [16](#)
 Cluster, [17](#)
 getDefaultCluster, [17](#)
 getID, [17](#)
 getNumberOfkThreads, [17](#)
 operator=, [17](#)

cluster_create
 cwrapper.h, [45](#)

cluster_destroy
 cwrapper.h, [45](#)

cluster_get_current
 cwrapper.h, [45](#)

cluster_get_default
 cwrapper.h, [45](#)

cluster_get_id
 cwrapper.h, [45](#)

cluster_get_number_of_kThreads
 cwrapper.h, [46](#)

condition_variable_create
 cwrapper.h, [46](#)

condition_variable_destroy
 cwrapper.h, [46](#)

condition_variable_empty
 cwrapper.h, [46](#)

condition_variable_signal
 cwrapper.h, [46](#)

condition_variable_signal_all
 cwrapper.h, [46](#)

condition_variable_wait
 cwrapper.h, [46](#)

ConditionVariable, [18](#)
 empty, [18](#)
 signal, [18](#)
 signalAll, [18](#)
 wait, [19](#)

connect
 Connection, [22](#)

Connection, [19](#)
 accept, [21](#)
 bind, [21](#)
 close, [21](#)
 connect, [22](#)
 Connection, [20](#), [21](#)
 getFd, [22](#)
 listen, [22](#)
 read, [22](#)
 recv, [22](#)
 recvfrom, [23](#)
 recvmsg, [23](#)
 recvmsg, [23](#)
 send, [24](#)
 sendmsg, [24](#)
 sendmsg, [24](#)
 sendto, [25](#)
 socket, [25](#)
 write, [25](#)

connection_accept
 cwrapper.h, [47](#)

connection_accept_connexion
 cwrapper.h, [47](#)

connection_bind
 cwrapper.h, [47](#)

connection_block_on_read
 cwrapper.h, [47](#)

connection_block_on_write
 cwrapper.h, [47](#)

connection_close
 cwrapper.h, [48](#)

connection_connect
 cwrapper.h, [48](#)

connection_create
 cwrapper.h, [48](#)

connection_create_socket
 cwrapper.h, [48](#)

connection_create_with_fd
 cwrapper.h, [48](#)

connection_destroy
 cwrapper.h, [49](#)

connection_get_fd
 cwrapper.h, [49](#)

connection_listen

- cwrapper.h, 49
- connection_read
 - cwrapper.h, 49
- connection_recv
 - cwrapper.h, 49
- connection_recvfrom
 - cwrapper.h, 49
- connection_recvmmsg
 - cwrapper.h, 50
- connection_recvmsg
 - cwrapper.h, 50
- connection_send
 - cwrapper.h, 50
- connection_sendmmsg
 - cwrapper.h, 51
- connection_sendmsg
 - cwrapper.h, 51
- connection_sendto
 - cwrapper.h, 51
- connection_socket
 - cwrapper.h, 52
- connection_write
 - cwrapper.h, 52
- create
 - uThread, 37
- currentUThread
 - uThread, 37
- currentkThread
 - kThread, 28
- cwrapper.h
 - cluster_create, 45
 - cluster_destroy, 45
 - cluster_get_current, 45
 - cluster_get_default, 45
 - cluster_get_id, 45
 - cluster_get_number_of_kThreads, 46
 - condition_variable_create, 46
 - condition_variable_destroy, 46
 - condition_variable_empty, 46
 - condition_variable_signal, 46
 - condition_variable_signal_all, 46
 - condition_variable_wait, 46
 - connection_accept, 47
 - connection_accept_connenction, 47
 - connection_bind, 47
 - connection_block_on_read, 47
 - connection_block_on_write, 47
 - connection_close, 48
 - connection_connect, 48
 - connection_create, 48
 - connection_create_socket, 48
 - connection_create_with_fd, 48
 - connection_destroy, 49
 - connection_get_fd, 49
 - connection_listen, 49
 - connection_read, 49
 - connection_recv, 49
 - connection_recvfrom, 49
 - connection_recvmmsg, 50
 - connection_recvmsg, 50
 - connection_send, 50
 - connection_sendmmsg, 51
 - connection_sendmsg, 51
 - connection_sendto, 51
 - connection_socket, 52
 - connection_write, 52
 - kThread_create, 52
 - kThread_destroy, 53
 - kThread_get_current, 53
 - kThread_get_current_pthread_id, 53
 - kThread_get_pthread_id, 53
 - kThread_get_total_number_of_kThreads, 53
 - mutex_acquire, 53
 - mutex_create, 53
 - mutex_destroy, 54
 - mutex_release, 54
 - ownerlock_acquire, 54
 - ownerlock_create, 54
 - ownerlock_destroy, 54
 - ownerlock_release, 54
 - semaphore_create, 54
 - semaphore_destroy, 54
 - semaphore_p, 54
 - semaphore_v, 54
 - uThread_create, 54
 - uThread_detach, 55
 - uThread_get_current, 55
 - uThread_get_id, 55
 - uThread_get_total_number_of_uThreads, 55
 - uThread_join, 55
 - uThread_migrate, 55
 - uThread_start, 56
 - uThread_terminate, 56
 - uThread_yield, 56
- empty
 - ConditionVariable, 18
- getCurrentCluster
 - uThread, 37
- getDefaultCluster
 - Cluster, 17
- getFd
 - Connection, 22
- getID
 - Cluster, 17
 - kThread, 28
 - uThread, 38
- getNumberOfkThreads
 - Cluster, 17
- getThreadNativeHandle
 - kThread, 28
- getTotalNumberOfkThreads
 - kThread, 29
- getTotalNumberOfUTs
 - uThread, 38

- IOPoller, [26](#)
- join
 - uThread, [38](#)
- kThread, [26](#)
 - currentkThread, [28](#)
 - getID, [28](#)
 - getThreadNativeHandle, [28](#)
 - getTotalNumberOfkThreads, [29](#)
 - kThread, [28](#)
 - operator=, [29](#)
- kThread_create
 - cwrapper.h, [52](#)
- kThread_destroy
 - cwrapper.h, [53](#)
- kThread_get_current
 - cwrapper.h, [53](#)
- kThread_get_current_pthread_id
 - cwrapper.h, [53](#)
- kThread_get_pthread_id
 - cwrapper.h, [53](#)
- kThread_get_total_number_of_kThreads
 - cwrapper.h, [53](#)
- listen
 - Connection, [22](#)
- migrate
 - uThread, [38](#)
- Mutex, [29](#)
 - acquire, [31](#)
- mutex_acquire
 - cwrapper.h, [53](#)
- mutex_create
 - cwrapper.h, [53](#)
- mutex_destroy
 - cwrapper.h, [54](#)
- mutex_release
 - cwrapper.h, [54](#)
- operator=
 - Cluster, [17](#)
 - kThread, [29](#)
 - uThread, [38](#)
- OwnerLock, [31](#)
 - acquire, [32](#)
 - release, [32](#)
- ownerlock_acquire
 - cwrapper.h, [54](#)
- ownerlock_create
 - cwrapper.h, [54](#)
- ownerlock_destroy
 - cwrapper.h, [54](#)
- ownerlock_release
 - cwrapper.h, [54](#)
- P
 - Semaphore, [34](#)
- pop
 - uThreadCache, [40](#)
- push
 - uThreadCache, [40](#)
- read
 - Connection, [22](#)
- recv
 - Connection, [22](#)
- recvfrom
 - Connection, [23](#)
- recvmsg
 - Connection, [23](#)
- recvmsg
 - Connection, [23](#)
- release
 - OwnerLock, [32](#)
- Semaphore, [33](#)
 - P, [34](#)
 - Semaphore, [33](#)
- semaphore_create
 - cwrapper.h, [54](#)
- semaphore_destroy
 - cwrapper.h, [54](#)
- semaphore_p
 - cwrapper.h, [54](#)
- semaphore_v
 - cwrapper.h, [54](#)
- send
 - Connection, [24](#)
- sendmsg
 - Connection, [24](#)
- sendmsg
 - Connection, [24](#)
- sendto
 - Connection, [25](#)
- signal
 - BlockingQueue, [14](#)
 - ConditionVariable, [18](#)
- signalAll
 - BlockingQueue, [15](#)
 - ConditionVariable, [18](#)
- socket
 - Connection, [25](#)
- start
 - uThread, [38](#)
- suspend
 - BlockingQueue, [15](#)
- terminate
 - uThread, [39](#)
- uThread, [34](#)
 - create, [37](#)
 - currentUThread, [37](#)
 - getCurrentCluster, [37](#)
 - getID, [38](#)
 - getTotalNumberOfUTs, [38](#)
 - join, [38](#)

- migrate, [38](#)
- operator=, [38](#)
- start, [38](#)
- terminate, [39](#)
- uThread, [37](#)
- yield, [39](#)
- uThread_create
 - cwrapper.h, [54](#)
- uThread_detach
 - cwrapper.h, [55](#)
- uThread_get_current
 - cwrapper.h, [55](#)
- uThread_get_id
 - cwrapper.h, [55](#)
- uThread_get_total_number_of_uThreads
 - cwrapper.h, [55](#)
- uThread_join
 - cwrapper.h, [55](#)
- uThread_migrate
 - cwrapper.h, [55](#)
- uThread_start
 - cwrapper.h, [56](#)
- uThread_terminate
 - cwrapper.h, [56](#)
- uThread_yield
 - cwrapper.h, [56](#)
- uThreadCache, [39](#)
 - pop, [40](#)
 - push, [40](#)
- wait
 - ConditionVariable, [19](#)
- write
 - Connection, [25](#)
- yield
 - uThread, [39](#)